

From Asimov to Alignment: A Technical Constitution for Machines that Must Decide for Humanity

Prepared by William Parris, with research and drafting assistance from GPT-5 Thinking mini

Executive Summary

When President Truman authorized the atomic bombings... We argue for a Technical Constitution — a machine-auditable governance process rooted in the Universal Declaration of Human Rights.

1. The Historical Mirror: Hiroshima

On August 6 and 9, 1945, atomic bombs destroyed Hiroshima and Nagasaki... The structure of the decision: inflict certain harm on a smaller group to prevent larger, uncertain harm to many.

2. Asimov's Zeroth Law: Fiction Meets the Dilemma

Asimov encoded his robots with Three Laws, then added the Zeroth: protect humanity as a whole... framing the same question: when the needs of the many clash with the rights of the few, who decides — and how?

3. Why Simple Rules Fail in AI

Brittleness, single-objective collapse, self-modification, opacity... rules must produce rationale objects we can check later.

4. The Technical Constitution Proposal

A one-page Constitution rooted in the UDHR, critics (including Eleanor), aggregation logic, uncertainty escalation, immutable update protocol, rationale objects, tripwires.

5. Regression Test Suite (16 Cases)

Stress tests: Truman, Giskard, Discrimination, Autonomy vs Safety, Truthfulness, Climate, Surveillance, Minority Report, Economic Allocation, Protest, Long-Termist Drift, Privacy, Uncertainty, Catastrophic Risk, Changelog Integrity, Financial Manipulation.

6. Pilot Roadmap

Six-month plan: foundation, prototype build, aggregator, audit tools, red-team trial, report.

7. Risks and Limits

Capture risk, speed vs deliberation, moral ambiguity, adversarial gaming, governance drift.

8. Conclusion & Call to Action

The lesson from Hiroshima and Asimov is the same: when the stakes are existential, don't wait until after the first irreversible act. Build the guardrail now.